



Kubernetes- Deepdive

Section 5 : Kubernetes deep-dive

PODs with YAML

Replication Controllers

Replication controller | Replica Sets

Hand-on for Replication controller | Replica sets

Replication controller

Replicaset

Scale & Test cluster

Scale up and down ReplicaSet

Test Replication-controller & Replica-sets

Deleting PODs created through Replicas

Deployment

Hand-on Deployment

Updates & Rollback

Rollout & Versioning

Deployment Strategy

Rollout update

Upgrades

Rollback

Session 6 : Kubernetes on cloud

Section 5 : Kubernetes deep-dive

PODs with YAML

- In this section will talk about creating the PODs using YAML based configuration file.
- We will be developing some YAML files which is going to create PODs for us in an automated fashion.
- Any file which you are going to create should be with .yml extension, For instance.

```
E0-first-pod.yml
```

- The kubernetes yml file should always have 4 top level fields (Mandatory).

```
apiVersion:  
kind:  
metadata:  
spec:
```

apiVersion :

- This is the version of the Kubernetes api we are using to create the objects.
- Depending on what we are trying to create we must use the right apiVersion.
- For now since we are working on PODs we will set the apiVersion as v1
- Other possible values for this field are

Kind	Version
POD	v1
Service	v1
ReplicaSet	apps/v1
Deployment	apps/v1

kind :

- Kind refers to the type of object we are trying to create, which in this case happens to be POD.
- Other possible values for this field would be [Service, Replicaset & Deployment]

metadata :

- Metadata is going to contain the details about the object [POD] you are going to create, Like its Name, Labels etc about your POD..
- Metadata mostly will be declared in form of Dictionary.
- Whatever you are going define under this field is considered as children of Metadata.

spec :

- Specification is where you are going to define the details about your container, Like name of your container, Image, port and volume details.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
    tire: frontend
spec:
  containers:
    - name: nginx-container
      image: nginx
```

Now your file is created, how to execute it ?

```
kubectl create -f E0-first-pod.yml
```

- **-f used to pass file**

```
kubectl apply -f E0-first-pod.yml
```

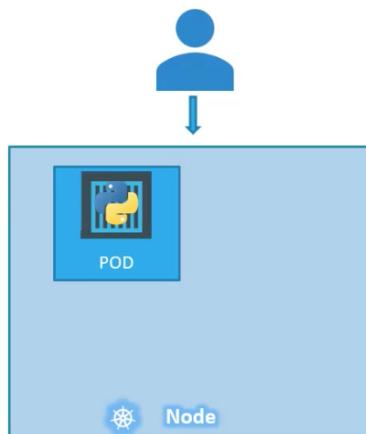
Replication Controllers

- Controllers are brain behind the Kubernetes.

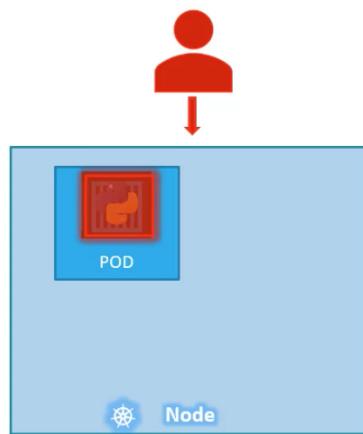
- They are the processes that monitor kubernetes objects and respond accordingly.
- In this section we are going to discuss specifically about Replication controller.

What is Replication controller ?

- Let's go back to our first scenario where we had single pod running our application.

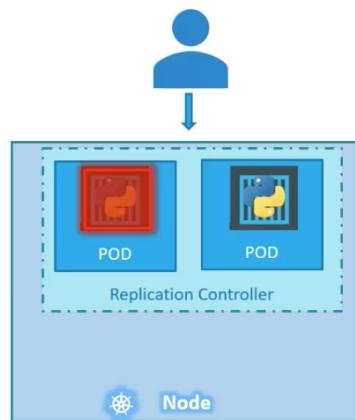


- What if for some reason our application got crashes & the Pod fails.? Users will no longer able to access the application hosted in Pod.

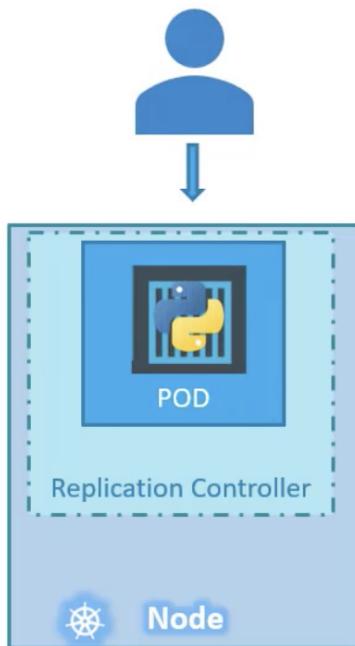


- To prevent users loosing access to application we would like to have more than one instance or Pods running at same time.
- In this way we can achieve the high availability of our application. If one Pod fails application will be still accessible from other Pod.

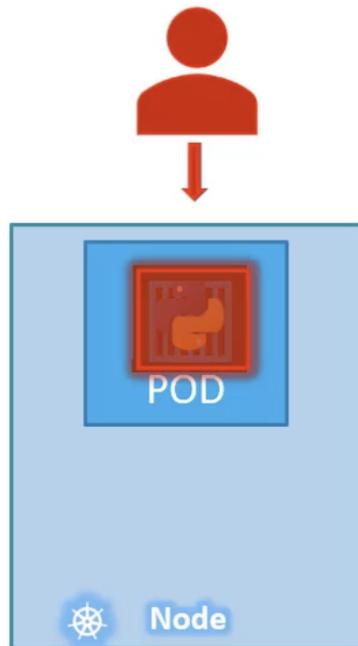
- **Replication controller helps us to run multiple instances of a single Pod as in cluster form to achieve High availability.**



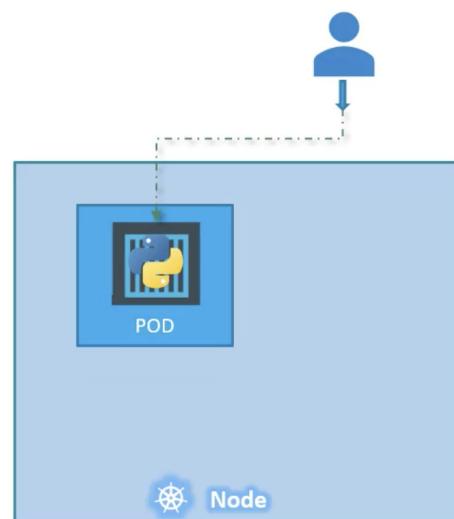
- **But does that mean you can't use replication factor if you plan to have single Pod ?**

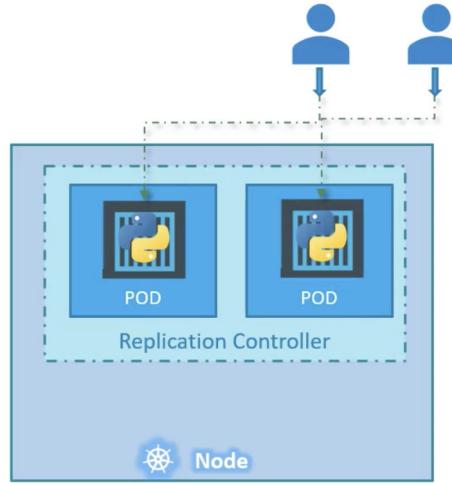


- **No "Even if you have single POD the replication factor can help by automatically bringing up the new POD when the existing one fails"**

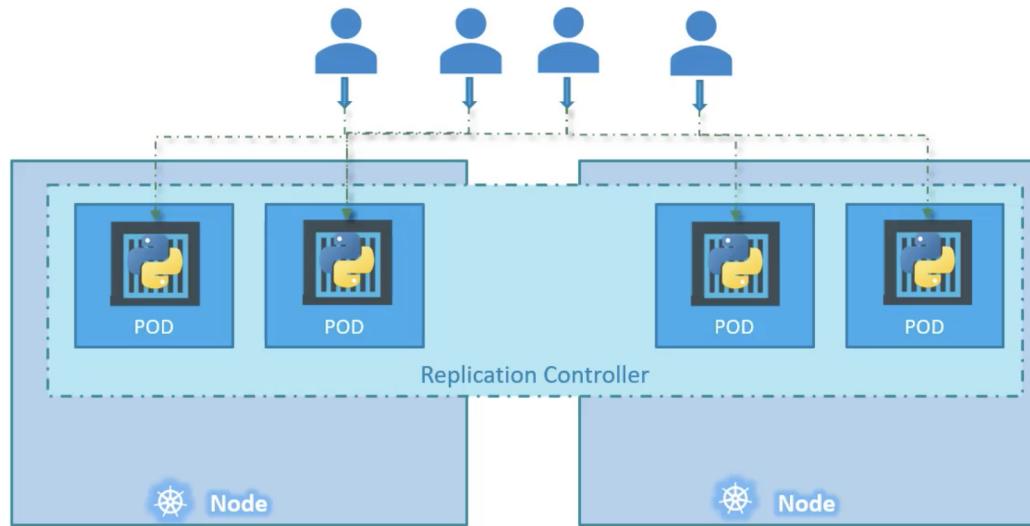


- Thus the replication controller ensures the specified number of PODs are running all time, Even if it is 1 or 100 PODs.
- Another reason to have replication controller is to create multiple PODs to share the Loads across them.
- Simple scenario which we discussed earlier, when the number of people accessing your application things should be fine.
- When the load increases your replication controller will automatically create required PODs to serve your customers.





- If the demand further increases and if we are running out of resources in first node we can deploy additional PODs across the other Nodes in the cluster.
- As you can see replication controller spans across multiple nodes in the cluster.
- This approach helps us to balance the load across multiple PODs on different nodes and scale our application when the demand increases.



Replication controller | Replica Sets

- It's important to understand that there are 2 similar terms Replication controller | Replica Set
- Both have the same purpose but they are not same.

- **Replication controller is the older technology that is been replaced by Replica Set.**
- **Replica set the new recommended way to setup replication.**
- **However whatever we have discussed in previous slides are applicable for both this functionalities.**
- **There are minor differences in the way how it works we shall look at those further.**

Hand-on for Replication controller | Replica sets

Replication controller

- Lets create an yaml file here for this example

```
E1-rc-definition.yml
```

- This file will have the default four fields as-usual [apiVersion, Kind, Metadata & spec]
- Under spec we are going to provide POD details under template section.

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: myapp-rc
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: nginx
      labels:
        app: nginx
        type: frontend
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
```

Execute file

```
kubectl apply -f E1-rcontroller.yml
```

To view the list of created replication controller

```
kubectl get replicationcontroller
```

To get more details about your replication controller

```
kubectl describe replicationcontroller myapp-rc
```

To see the PODs

```
kubectl get pods
```

Replicaset

Let's talk about Replicaset now

- The code is going to look similar here but with some changes and additional parameters.
- Let's create new file here.

```
E2-replicaset.yml
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-prod
      labels:
        app: myapp
        type: front-end
    spec:
```

```
  containers:
    - name: nginx-container
      image: nginx
  replicas: 3
  selector:
    matchLabels:
      type: front-end
```

- Some changes we did here is..
- Change of apiVersion [Ensure if you are making mistake in apiVersion you might likely to get error like this]

```
error: unable to recognize "replicaset-definition.yml": no matches for /, Kind=ReplicaSet
```

- Adding additional section called "selector" under replicas.
- The "selector" section helps replicaset to identify what pods falls under it.
- But why would you have to specify what PODs falls under it ? As you can see we have already provided the contents of the PODs in template section already.
- The reason for doing it is because "replicaset can also manage PODs that were not created as part of the replicaset creation"
- For Example : Lets assume there were some PODs created before the creation of Replicaset's creation.
- Now the Replicaset will also consider those PODs while creation if incase the label specified in the selector section matches that independent POD.
- The major difference between ReplicationController & ReplicaSet is this "selector section"

Execute file

```
kubectl create -f E2-replicaset.yml
```

To view the list of created replication controller

```
kubectl get replicaset
```

To get more details about your replicaset

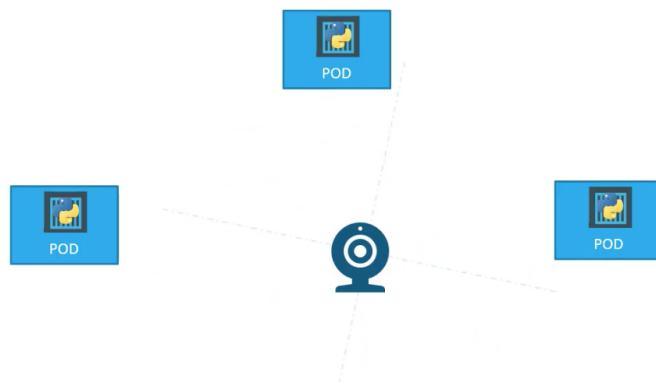
```
kubectl describe replicaset myapp-replicaset
```

To see the PODs

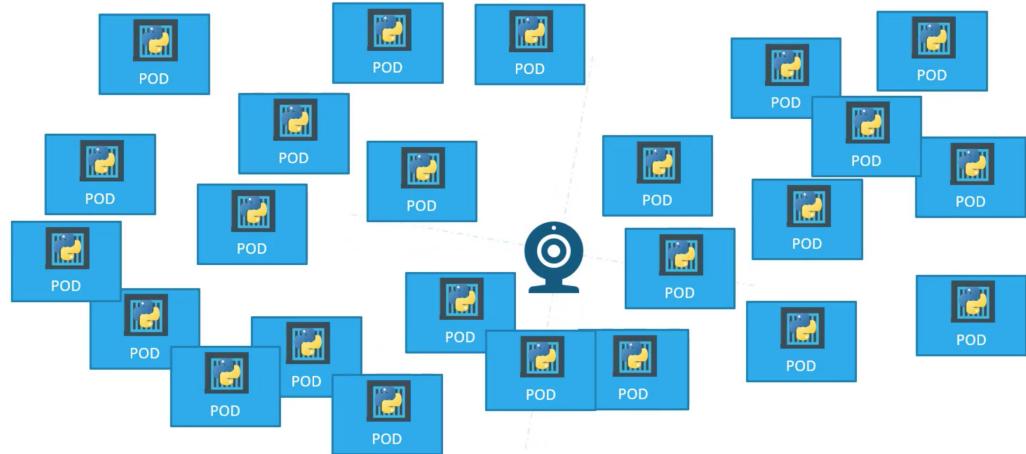
```
kubectl get pods
```

Labels and Selectors

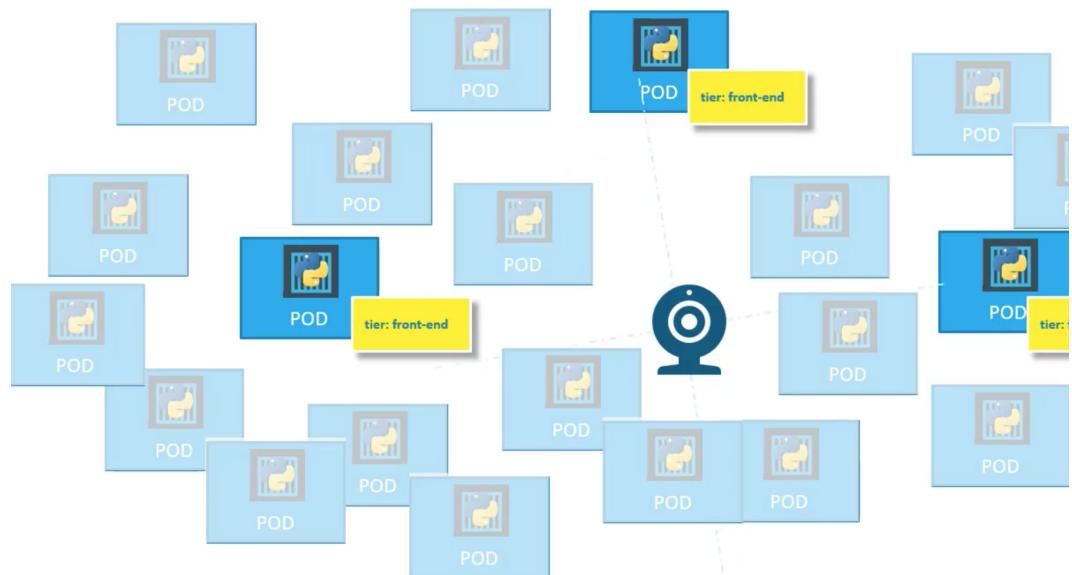
- **Whats the deal with Labels and Selectors ? and why do we label PODs and objects in kubernetes.?**
- **Let's look into a simple scenario.**
 - **Just assume we have 3 PODs with instances configured in it. We would like to create Replication controller or Replicaset to ensure that we have 3 active PODs at any time and all the time.**



- **As you already know the role of Replication Controller or ReplicaSet is to monitor the PODs and recreate if any of those fails.**
- **Now how does the replicaset know which PODs to be monitored ? because we might have many PODs running in your cluster.**



- This is where labels are going to be handy to filter the PODs which need to be monitored by ReplicaSet.
- In this way Labels might be handy in many places across Kubernetes cluster.



Scale & Test cluster

Scale up and down ReplicaSet

- Let's see how to scale the ReplicaSet.
- For now we started with 3 Replicas, what if we want to scale it in future.
- Lets assume we need to scale our Replicas from 3 to 6.
- There are multiple ways to do this

Step 1 : Update the code with ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-prod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 6
  selector:
    matchLabels:
      type: front-end
```

Then run the code like below

```
kubectl replace -f E2-replicaset.yml
```

Step 2 : Update the code with scale command

```
kubectl scale --replicas=6 -f E2-replicaset.yml
```

Similarly to scale down the cluster

```
kubectl scale --replicas=2 -f E2-replicaset.yml
```

Step 3 : Changing the replicas in running configuration

- We can also change this configuration at running config file of replicas.

- When we open this file we can see the running configuration of replicaset.
- Note that this is not the actual file which we created, This file is created by Kubernetes in memory.
- Kubernetes allows us to edit the objects from this file as well.
- There might be lot of additional fields available which we have not specified in our YAML file and there are generated by Kubernetes itself.

```
kubectl edit replicaset myapp-replicaset
```

Test Replication-controller & Replica-sets

Now lets try deleting the PODs and just see if its getting created.

```
kubectl delete pods myapp-replicaset-<pod-name>
```

Eventually POD should be recreated now.

Deleting PODs created through Replicas

How to delete the machines created through Replicas & Replication-controller.

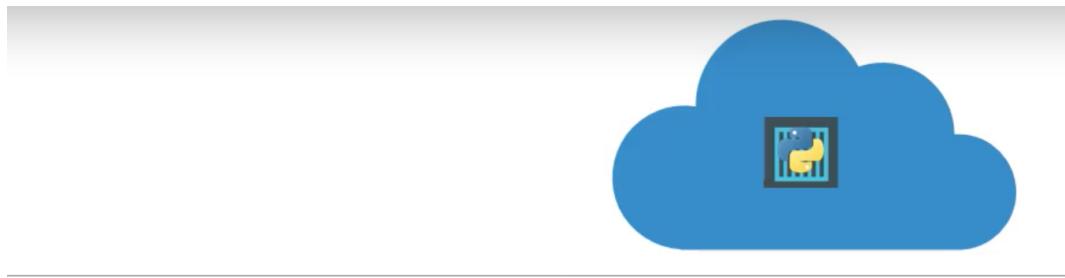
```
kubectl delete replicationcontroller myapp-rc
```

```
kubectl delete replicaset myapp-replicaset
```

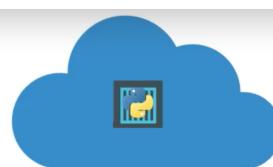
Deployment

Let's talk about concept called Deployment.

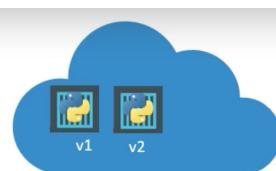
- Let's first think about how we might want to deploy the applications in Production environment to understand Deployments.
- Say for example you have web server which needs to be deployed in Production environment.



- For sure you are not going to stop having one server for serving your customers, eventually over the period number of servers might increase.

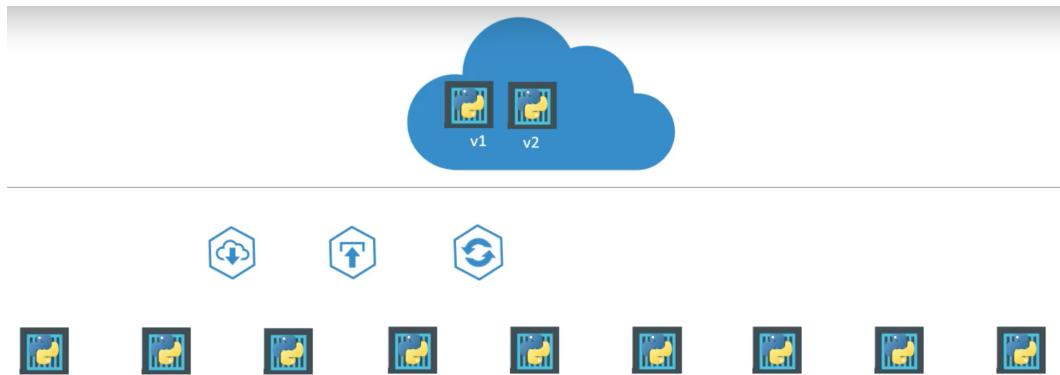


- Secondly, whenever newer versions of application become available in Docker registry. You may like to upgrade your applications seamlessly.



- However, while you upgrade your instances you do not want them to upgrade all of them once because it may impact users accessing the

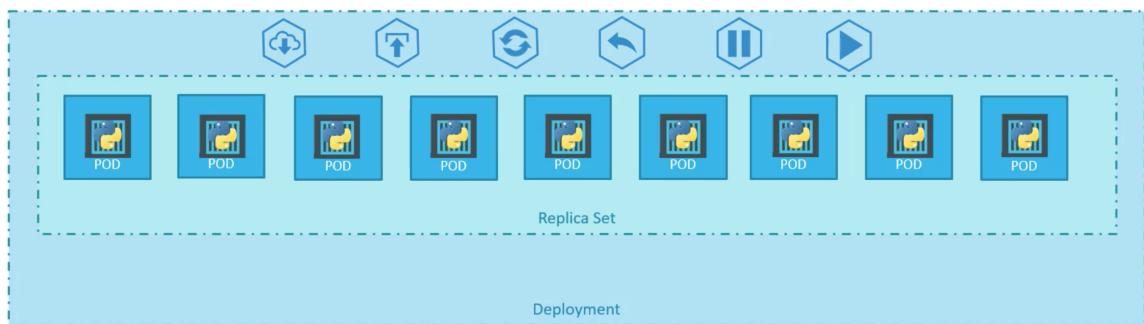
application. Hence you might want to upgrade the instance one after the other which is known as Rolling updates.



- Suppose one of the upgrades which you perform encountered an unexpected error and you are asked to undo the recent upgrades which is known as Rollback changes.



- So far in this training we have discussed about PODs which helps to encapsulate your Instances.
- Multiple such PODs are deployed using Replica Sets.
- Then comes the kubernetes object which comes higher in the hierarchy.



- The main objective of Deployment is to help you in Making multiple changes in your environment such as
 - upgrading the underlying instances application versions seamlessly.
 - Rollback if incase of any problems encountered.

Hand-on Deployment

Let's see how to create deployments.

- Creating file

```
E3-deployment.yml
```

- Changes here are kind & name of object.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deployment
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-prod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      type: front-end
```

Execute code

```
kubectl apply -f E3-deployment.yml
```

Verify all the details

```
kubectl get deployment
```

```
kubectl describe deployment myapp-deployment
```

```
kubectl get replicaset
```

```
kubectl get pods
```

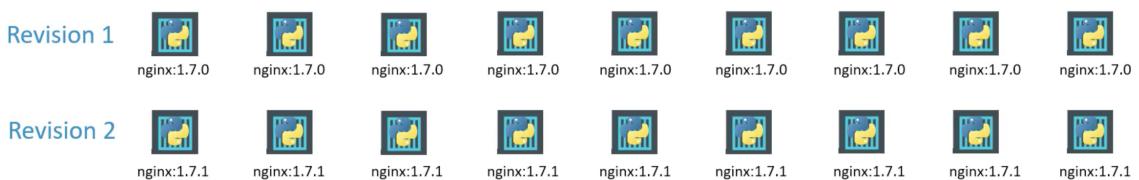
To see all these details in one command

```
kubectl get all
```

Updates & Rollback

Rollout & Versioning

- When you first create deployment it triggers rollout.
- New rollout creates new deployment revision.
- In future when application is upgraded a new rollout is triggered and it creates new deployment revision.
- This approach actually helps us to keep track of the changes made to our deployments & enables us to rollback to previous version when and where required.



How to check the rollout and revision status ?

```
kubectl rollout status deployment/myapp-deployment
```

```
kubectl rollout history deployment/myapp-deployment
```

Deployment Strategy

There are two different deployment strategy.

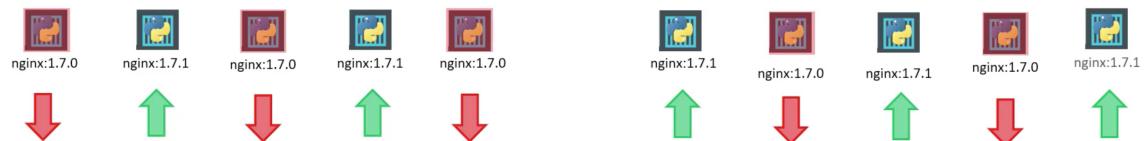
Strategy 1 : Recreate

- For example lets assume you have 5 web application instance running.
- One way to upgrade all these instance to newer version is "destroy all the older one and recreate new instance with newer version"
- But the problem with this approach is application might be down and inaccessible for your customers until new instances are coming up.



Strategy 2 : Rolling update

- Through this strategy we are going to bring down the instance one by one and replace it with newer version before we move to the next one.
- Through this approach you can achieve zero downtime.
- Note : This is the default approach followed by kubernetes even if you are not specifying anything in your YAML configuration file.



Rollout update

Let's try to rollout the update for existing containers running.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deployment
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-prod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx:1.7.1
replicas: 3
selector:
  matchLabels:
    type: front-end
```

Apply code

```
kubectl apply -f E4-deploy-update.yml
```

You can also do this upgrade in other way without updating your YAML file.

```
kubectl set image deployment/myapp-deployment nginx=nginx:1.9.1
```

Once done you can check the new revision using command

```
kubectl rollout history deployment/myapp-deployment
```

Now you can understand the difference between recreate and rollout.

- Recreate will zero down all the instance first before create new.

- Whereas rollout method is going to do it one by one.

```
C:\Kubernetes>kubectl describe deployment myapp-deployment
Name:           myapp-deployment
Namespace:      default
CreationTimestamp: Sat, 03 Mar 2018 17:01:55 +0800
Labels:         app=myapp
                type=front-end
Annotations:   deployment.kubernetes.io/revision=2
                kubectl.kubernetes.io/last-applied-configuration={"apiVersion":"apps/v1","kind":"Deployment","metadata":{"name": "myapp-deployment","namespace": "default","labels": {"app": "myapp","type": "front-end"}, "annotations": {"deployment.kubernetes.io/revision": "2"}}, "status": {"availableReplicas": 5, "desiredReplicas": 5, "unavailableReplicas": 0, "conditions": [{"type": "Available", "status": "True", "reason": "MinimumReplicasAvailable"}, {"type": "Progressing", "status": "True", "reason": "NewReplicaSetAvailable"}], "replicaSets": [{"name": "myapp-deployment-54c7d6ccc", "status": "Available", "replicas": 5}], "events": [{"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled up replica set myapp-deployment-6795844db8 to 5"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled down replica set myapp-deployment-6795844db8 to 0"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "56s", "from": "deployment-controller", "message": "Scaled up replica set myapp-deployment-54c7d6ccc to 5"}]}

C:\Kubernetes>kubectl describe deployment myapp-deployment
Name:           myapp-deployment
Namespace:      default
CreationTimestamp: Sat, 03 Mar 2018 17:16:53 +0800
Labels:         app=myapp
                type=front-end
Annotations:   deployment.kubernetes.io/revision=2
                kubectl.kubernetes.io/last-applied-configuration={"apiVersion":"apps/v1","kind":"Deployment","metadata":{"name": "myapp-deployment","namespace": "default","labels": {"app": "myapp","type": "front-end"}, "annotations": {"deployment.kubernetes.io/revision": "2"}}, "status": {"availableReplicas": 5, "desiredReplicas": 5, "unavailableReplicas": 0, "conditions": [{"type": "Available", "status": "True", "reason": "MinimumReplicasAvailable"}, {"type": "Progressing", "status": "True", "reason": "ReplicaSetUpdated"}], "replicaSets": [{"name": "myapp-deployment-67749c58c", "status": "Available", "replicas": 1}, {"name": "myapp-deployment-7057db08d", "status": "Available", "replicas": 5}], "events": [{"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled up replica set myapp-deployment-67749c58c to 5"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled down replica set myapp-deployment-67749c58c to 1"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled up replica set myapp-deployment-7057db08d to 3"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled down replica set myapp-deployment-7057db08d to 2"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled up replica set myapp-deployment-7057db08d to 4"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled down replica set myapp-deployment-7057db08d to 2"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled up replica set myapp-deployment-7057db08d to 5"}, {"type": "Normal", "reason": "ScalingReplicaSet", "age": "1m", "from": "deployment-controller", "message": "Scaled down replica set myapp-deployment-7057db08d to 1"}]}
```

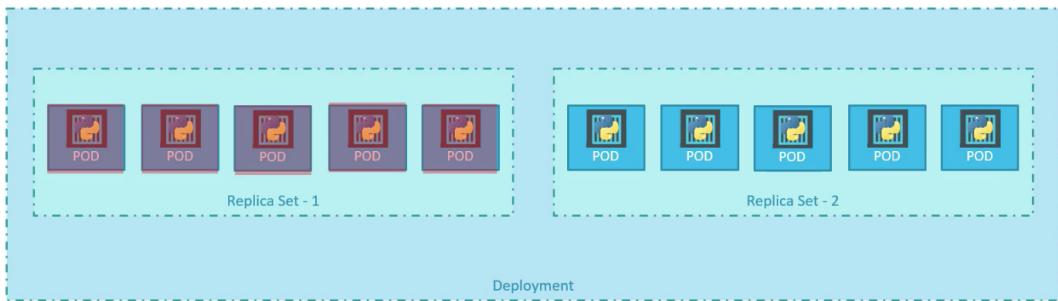
Upgrades

Let's all try to understand behind the woods when you are performing all these steps.

- When you rollout first deployment your instances will be running under Revision 1



- When you upgrade your instances kubernetes will create new Revision with updated instances & will bring down all the instance running in current Revision.



This can be verified using command

```
kubectl get replicaset
```

Rollback

Now how to rollback from currnt updated revision to older revision if incase you are encountering some problem with your application after upgrade.

```
kubectl rollout undo deployment/myapp-deployment
```

This command wiill rollback from revision2 to revision 1 again, You can find this from command.

```
kubectl get replicaset
```

Session 6 : Kubernetes on cloud

You can also create kubernetes cluster in cloud.

- In AWS kubernetes cluster is refered as Elastic Kubernetes Service (EKS)
- In Azure kubernetes cluster is refered as Azure Kubernetes Service (AKS)
- In GCP kubernetes cluster is refered as Google Kuberenets Engine (GKE)

Prerequisites to setup kubernets cluster in AWS.

- AWS account
- Knows some AWS basics

